


Applied mathematics in Deep Learning


Picture to numbers

What is the input format in a deep learning network?

Grayscale image = vector = $H \times W$



RGB image = vector = $C \times H \times W$



Picture to numbers

What is the input format in a deep learning network?

$$X \in \mathbb{R}^n$$

Where

- **X** contains the pixel values of an image
- \in means it's an "Element of"
- \mathbb{R}^n means "a vector with n real numbers."

So we say: "**X** is a vector with **n** numbers representing pixel values."

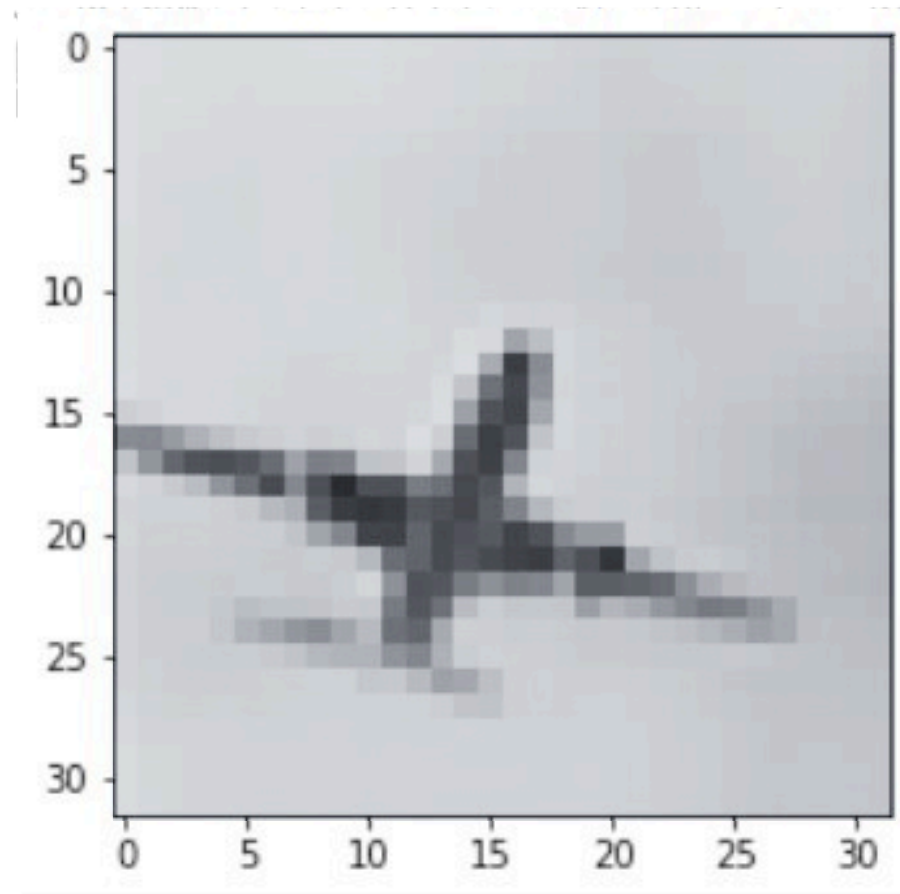
Picture to numbers

What is the input format in a deep learning network?

Example 32x32 image


For **grayscale**

$32 \times 32 = 1024$ pixels = $x \in \mathbb{R}^{1024}$



For **RGB**

$32 \times 32 \times 3 = 3072$ pixels = $x \in \mathbb{R}^{3072}$



Weights & Biases

How the network learns

Weights (W) - How important is each input?


Biases (B) - How important is it to respond to an input

For a 'forward pass' in a network for a single input:

$$\text{Output} = f((w_0 * x) + (w_1 * b))$$


Where

- w_0 controls how important the input is
- w_1 shifts the output results
- f is the activation function



Layered architecture

Input, output and hidden layers




Activation functions


Add non linearity

Activation function: decides the output of the neuron


Tanh




ReLU



Sigmoid



Linear




Activation functions: Linear functions


A linear function means:

- If you double the input, the output doubles too
- If you add two inputs, they are added too the output too

$$f(x) = x + 0$$



$$f(x) = 3x + 5$$





Activation functions: ReLU

ReLU (Rectified Linear Unit): helps negative neurons 'turn off'

$$f(x) = \max(0, x)$$

Where:

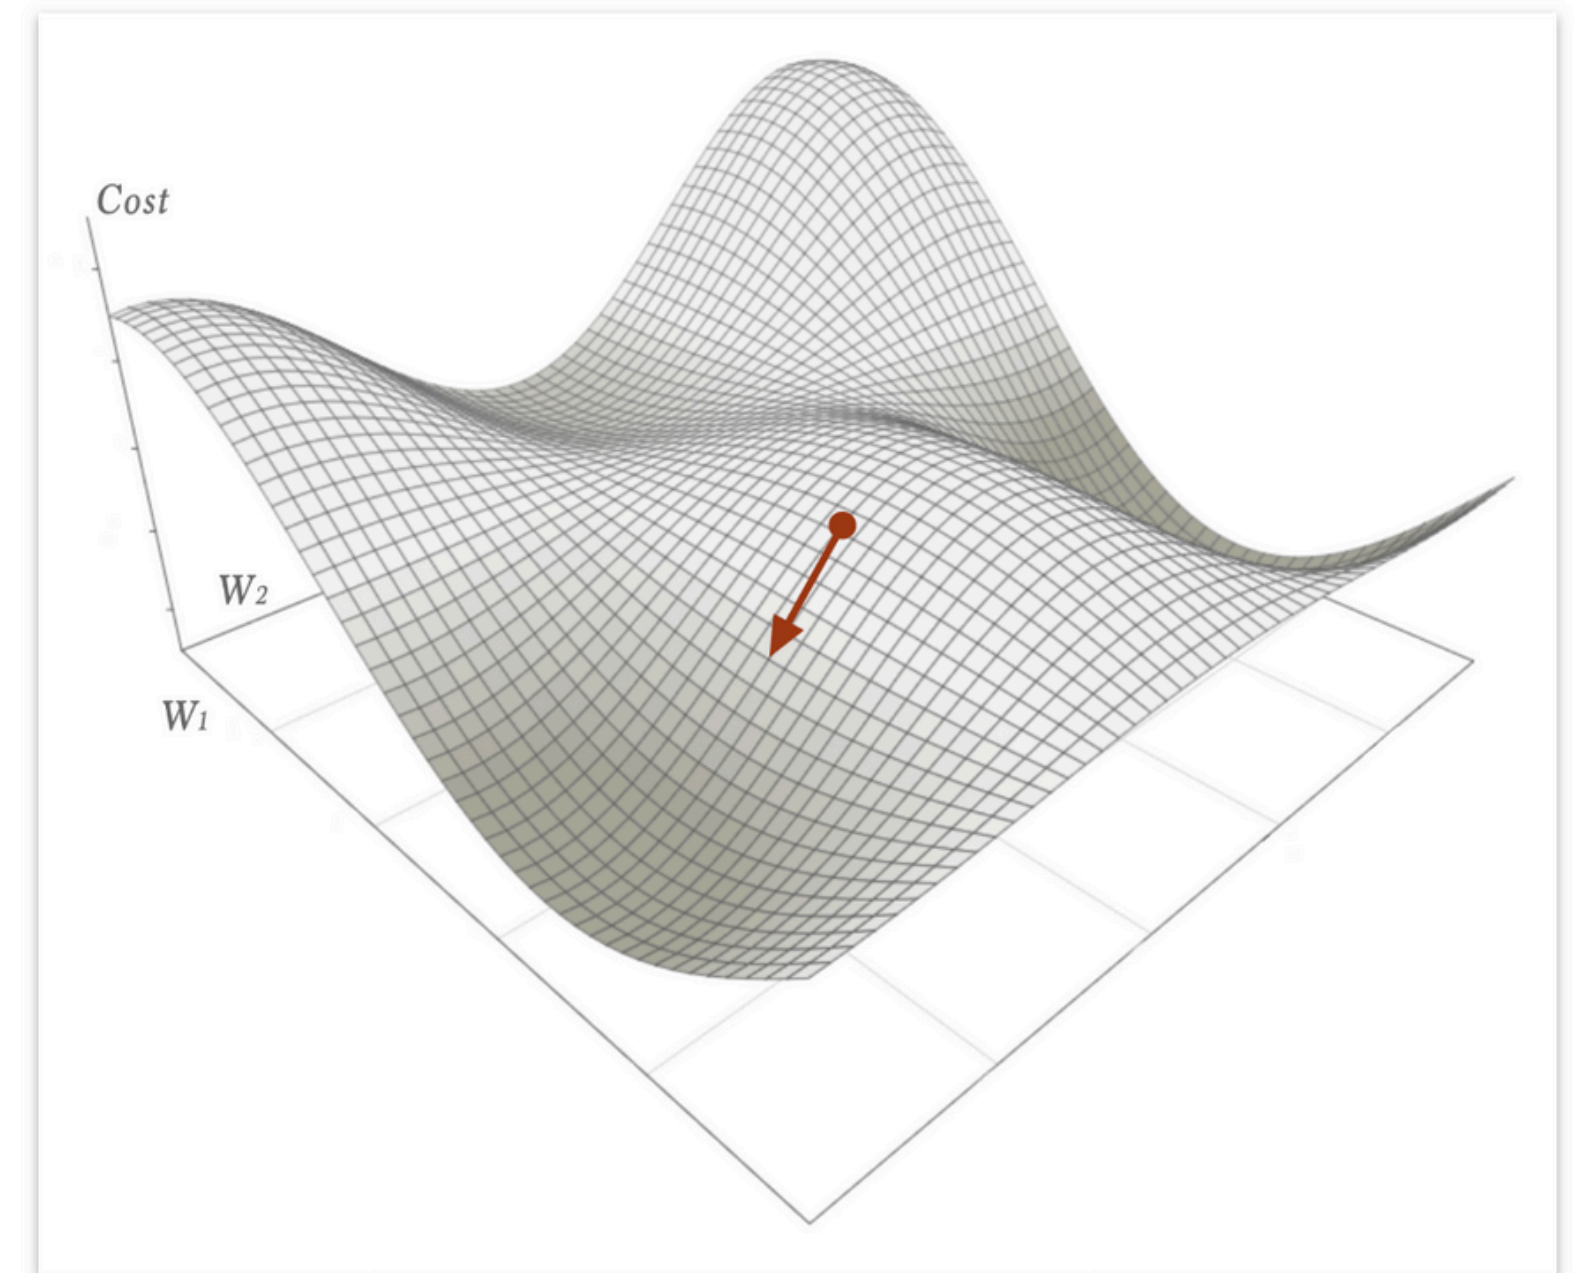
- For $x < 0$: the output is always 0 \rightarrow flat line
- For $x \geq 0$: the output equals $x \rightarrow$ diagonal line going up



Loss landscape

“How much is the network wrong per step?”

- If the model is confident and correct, the loss is **small**.
- If the model is wrong or unsure, the loss is **big**.



Loss landscape: Cross Entropy (CE)

Used when classes are balanced

$$CE_i = -\log(\text{softmax}(z_i)[y_i])$$

Where

- z_i : the model outputs (numbers)
- y_i : the correct class (index)
- $\text{softmax}(z_i)[y_i]$: the predicted probability for the correct class

Loss landscape: Focal Loss

$$FL_i = \alpha(1 - p_i)^\gamma \cdot (-\log(\text{softmax}(z_i)[y_i]))$$

Where

- **α** balances the contribution between positive and negative classes.
- **γ** controls the value of weights for harder to classify examples
- **p** the predicted probability

Making a prediction: Logits

3 classes: horse, dog, bird.

Final layer output: [2.0, 1.0, 0.1]



[2.0, 1.0, 0.1]
↓ ↓ ↓
[horse, dog, bird]

Making a prediction: Softmax

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\text{sum of } e^{\text{all } x_j}}$$

Where

- x_i the score of the class i (probability)
- x_j the scores (logits) of all classes
- e^{x_i} the exponential (raising a number to a power) for class i

